

Challenges and solutions when handling imbalanced datasets

Week 7 – Advanced Topic 5


Outlines

- Introduction
- Challenges
- Solutions
- Summary


Imbalanced dataset

A dataset with unequal representation of classes.

Class Distribution in an Imbalanced Two-Class Dataset



Class Distribution in an Imbalanced Multi-Class Dataset



Classification

- For **classification** models the most commonly used metric is **accuracy** (*proportion of **correctly classified** classes by the model*)
- Accuracy works** only if the dataset is **balanced** (*a **similar number** of samples in a dataset represents **each class***)

How to detect:

- use **Cohen's kappa coefficient** (k) or **Matthews correlation Coefficient** (MCC) instead of accuracy (*or at least, use **not only accuracy***)

Example:

F1	F2	F3	F4	F5	Class
2	4	2	8	5	A
1	3	1	7	2	A
5	5	1	1	7	A
3	7	1	2	6	A
4	3	2	6	3	A
5	6	3	5	3	B
3	6	1	7	4	B
1	5	1	6	2	B
1	4	3	2	6	B
6	7	2	8	4	B

Always
predict
class A

Predicted	Correct?
A	TRUE
A	TRUE
A	TRUE
A	TRUE
A	TRUE
A	FALSE
A	FALSE
A	FALSE
A	FALSE
A	FALSE

Number of correct classifications

$$\text{Accuracy} = \frac{5}{10} = 50\%$$

Total number of classifications

F1	F2	F3	F4	F5	Class
2	4	2	8	5	A
1	3	1	7	2	A
5	5	1	1	7	A
3	7	1	2	6	A
4	3	2	6	3	A
5	6	3	5	3	A
3	6	1	7	4	A
1	5	1	6	2	A
1	4	3	2	6	A
6	7	2	8	4	B

Always
predict
class A

Predicted	Correct?
A	TRUE
A	TRUE
A	TRUE
A	TRUE
A	TRUE
A	TRUE
A	TRUE
A	TRUE
A	TRUE
A	FALSE

Number of correct classifications

$$\text{Accuracy} = \frac{9}{10} = 90\%$$

Total number of classifications

Aleksei

Source: [2108.02497] How to avoid machine learning pitfalls: a guide for academic researchers (arxiv.org)

Imbalance in classification – we know how to detect it, but how to mitigate it? *[ChatGPT](#) on the topic*

Case: Out of 10,000 transactions, only 200 are fraudulent (class 1), and 9,800 are non-fraudulent (class 0). The dataset is **highly imbalanced**, with 2% fraudulent and 98% non-fraudulent transactions.

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score

# X: Features, y: Target (0 for non-fraud, 1 for fraud)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Initialize Logistic Regression without class balancing
model = LogisticRegression()

# Train the model
model.fit(X_train, y_train)

# Predictions
y_pred = model.predict(X_test)

# Evaluate the model
print(classification_report(y_test, y_pred))
```

Aleksei

	precision	recall
0	0.98	1.00
1	0.50	0.02

[Precision and recall - Wikipedia](#)

[ChatGPT](#) on Precision and Recall

$$\text{Precision} = \frac{\text{True Positives (TP)}}{\text{True Positives (TP)} + \text{False Positives (FP)}}$$

"Of all the instances that the model predicted as positive, how many were actually positive?"

$$\text{Recall} = \frac{\text{True Positives (TP)}}{\text{True Positives (TP)} + \text{False Negatives (FN)}}$$

"Of all the actual positive instances, how many did the model correctly identify?"

Imbalance in classification – we know how to detect it, but how to mitigate it? *ChatGPT on the topic*

Aleksei

```
# Initialize Logistic Regression with class balancing
model_balanced = LogisticRegression(class_weight='balanced')

# Train the model
model_balanced.fit(X_train, y_train)

# Predictions
y_pred_balanced = model_balanced.predict(X_test)

# Evaluate the model
print(classification_report(y_test, y_pred_balanced))
```

How the weight is used:

- The **logistic loss** function is **modified** by multiplying the **loss associated with each sample** by its **class weight**. If a class has a **higher weight**, the **errors associated with that class** will have a **greater impact on the optimization process**, pushing the model to reduce misclassifications for that class.

Note that these weights will be multiplied with `sample_weight` (passed through the fit method) if `sample_weight` is specified.

	precision	recall
0	0.98	1.00
1	0.50	0.02

Training **without** class balancing



	precision	recall
0	0.99	0.97
1	0.30	0.73

Training **with** class balancing

[LogisticRegression — scikit-learn 1.5.2 documentation](#)

Solution: balance classes if they are imbalanced, **use robust metrics** to reveal the imbalance

Solutions

- Use of robust metrics
- Apply cost-sensitive functions
- Resampling of classes
- Algorithmic solutions

Classification – robust metrics

- **Cohen's Kappa (k)**: This metric considers the chance of agreement between the model and the data occurring by random guessing. It ranges from **-1 (complete disagreement)** to **+1 (perfect agreement)**, where **0 means no agreement beyond chance**.

[Cohen's kappa – Wikipedia](#)

- **Matthews Correlation Coefficient (MCC)**: MCC is a **balanced** measure that considers **all four values from the confusion matrix (true positives, true negatives, false positives, and false negatives)**. It produces a value between -1 and 1, where **1 is a perfect prediction**, **0 is a random prediction**, and **-1 indicates total disagreement**.

[Phi coefficient – Wikipedia](#)

- Or simply use **Precision and Recall** together as metrics if you prefer simpler metrics

Cost-sensitive learning

- Defining a **cost matrix** that specifies different penalties for each type of misclassification
- The cost matrix is designed to specify **higher** penalties for errors in the **minority** class.

1		Actual Negative	Actual Positive
2	Predicted Negative	$C(0,0)$, TN	$C(0,1)$, FN
3	Predicted Positive	$C(1,0)$, FP	$C(1,1)$, TP




1		Actual Negative	Actual Positive
2	Predicted Negative	0	88
3	Predicted Positive	5	0

	C_1	C_2	C_3
\hat{C}_1	0	0.5	1
\hat{C}_2	2	0	1
\hat{C}_3	3	0.5	0


Resampling

- Aims to rebalance the class distribution by
 - Undersampling: **Remove** observations from the majority class
 - Oversampling: **Duplicate** observations from the minority class

Undersampling




Oversampling



Undersampling Techniques (1/4)

- **Random undersampling:** Randomly select and discard samples from the majority class until the desired class balance is achieved



Undersampling Techniques (2/4)


- **Cluster-based undersampling:** Uses clustering algorithms to group similar samples from the majority class and then select representative samples from each cluster



Original Dataset



Calculating Centroids




Selecting Farthest Points



Resampled Dataset

Undersampling Techniques (3/4)


- **Tomek links undersampling:** Removes noisy and borderline samples from the majority class




Undersampling Techniques (4/4)

- **NearMiss undersampling:** Uses the distance between the samples to select the ones to keep or discard from the majority class


Near Miss



Original Dataset




Selecting Samples



Resampled Dataset

Oversampling Techniques (1/2)


- **Random:** Selects the samples randomly and generate new samples in minority class



Oversampling Techniques (2/2)

- **Synthetic Minority Oversampling Technique (SMOTE)**

Finds 'k' nearest neighbour data points of minority class and creates synthetic data points on the lines joining the primary point



Drawbacks of resampling

- **Overfitting** towards minority class
- Chances of adding **noise** and **overlap**
- Model fine-tuned to balanced dataset might fail in dynamic environment with **changing data distributions**
- Sequential data, such as time series, loses **temporal** information when resampled

Solutions – Trying different algorithms

Classification Trees

Nodes are split based on *Impurity* or *Entropy*, which automatically consider class distributions at each split.




Classification trees can *focus on the minority class* in regions of the feature space where it is dominant, even if the overall dataset is imbalanced.

Solutions – Trying different algorithms

Anomaly Detection Approach

- Treat the minority class as anomalies or outliers.
- Any observation that does not fit the majority class is considered an anomaly or outlier.
- Apply anomaly detection techniques such as **One-Class SVM**, **Isolation Forest**, or **Autoencoders**.



Anomaly Detection image from [Wikipedia](#) (CC BY-SA 4.0)

Summary

- Imbalance data can cause **overfitting** of the majority class
- Use **resampling** methods for having a balanced dataset
- Use **cost sensitive** learning \ class weighting for making the model focus on the minority class
- Handling class imbalance is best suited for task like **classification** but not for anomaly detection