



# Naïve Bayes Classifier

---

Group 1:

Thale Eliassen Fink, Nicole Quattrini, Cameron Louis Penne, Aria Alinejad, Anette Fagerheim Bjerke, M. Tsaqif Wismadi, Bjørnar Kaarevik



# Outline

- Bayes' Theorem
- Naïve Bayes Classifiers
- Different types
- Examples
- Common cases
- Pros & Cons
- What not to do


# Bayes' Theorem

- A type of conditional probability of an event:

$$P(A|B) = \frac{P(B|A) \cdot P(A)}{P(B)}$$

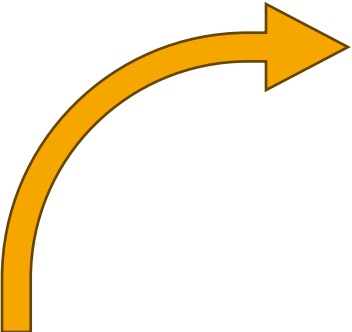
- $P(A|B)$  : probability of event  $A$  occurring given that  $B$  is true;
- $P(B|A)$  : probability of event  $B$  occurring given  $A$  is true;
- $P(A)$  &  $P(B)$  : probabilities of  $A$  and  $B$  respectively.

# Bayes' Theorem - Example

- The probability of having Really Terrible Disease  $A = 0.01$  ()
- The probability of *testing positive* given you have the disease  $= 0.99$

**Q:** What is the probability of having the disease given the positive test result?

# Bayes' Theorem - Example

- 
- $P(D)$  = Prob. having disease = 0.01
  - $P(D')$  = Prob. not having disease = 0.99
  - $P(T|D)$  = Prob. +ve test *given* having disease = 0.99
  - $P(T|D')$  = Prob. +ve test *given* not having disease = 0.01
  - $P(D|T)$  = ??

- The probability of having Really Terrible Disease A = 0.01 (🤖)
- The probability of *testing positive* given you have the disease = 0.99

**Q:** What is the probability of having the disease given the positive test result?

# Bayes' Theorem - Example

- $P(D)$  = Prob. having disease = 0.01
- $P(D')$  = Prob. not having disease = 0.99
- $P(T|D)$  = Prob. +ve test *given* having disease = 0.99
- $P(T|D')$  = Prob. +ve test *given not* having disease = 0.01
- $P(D|T)$  = ??

$P(T)$  = Total probability of testing positive  
= **0.0198**

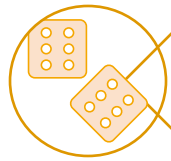
$$P(D|T) = \frac{P(T|D) \cdot P(D)}{P(T)}$$

$$P(D|T) = \frac{0.99 \cdot 0.01}{0.0198}$$

$$P(D|T) \approx 0.5$$

Even with a positive test result, there is a 50% chance you have the disease.

# Naïve Bayes Classifiers



Family of classifiers based on Bayes' Theorem



**Naïve**: assume feature independence



Used for classification

# Naïve Bayes Classifiers

## **Assumptions:**

- Feature independence
- Normal distribution (continuous variables)
- Multinomial distribution (categorical variables)
- Equal importance across features
- No missing data



# Naïve Bayes Classifiers



$$P(Y|X) = \frac{P(X|Y) \cdot P(Y)}{P(X)}$$

$Y$  - class

$X$  - feature

$P(Y = y_1)$  = Number of  $y_1$  in data set / Total elements in  $Y$

# Multinomial & Bernoulli Naive Bayes

## Bernoulli

$$P(A|\text{yellow doc.}) = \frac{\# \text{yellow docs. with feature}}{\# \text{yellow docs.}} = \frac{2}{3}$$

## Multinomial

$$P(A|\text{yellow doc.}) = \prod \frac{\#A \text{ in this yellow doc.}}{\# \text{features in this yellow doc.}} = \frac{3}{5} * \frac{0}{3} * \frac{4}{4} = \frac{0}{60}$$

ABCAA

BBC

AAAA



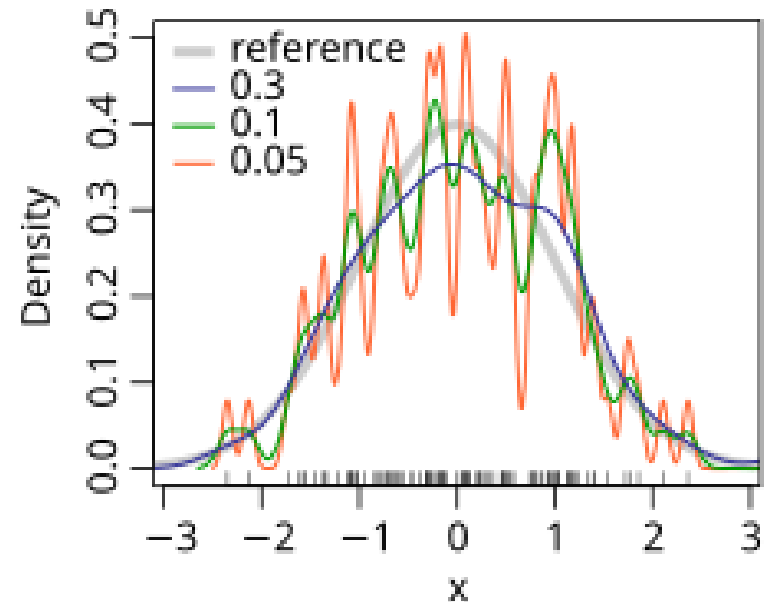
# Gaussian Naïve Bayes

- Assumes normal distributed features
  - Often far from true
- Handles continuous data
  - Binning or Kernel density estimation can also be used

$$P(x_i|y_j) = \frac{1}{\sqrt{2\pi\sigma_j^2}} \exp\left(-\frac{(x_i - \mu_j)^2}{2\sigma_j^2}\right)$$

# Kernel density estimation (KDE)

- Non-parametric method to estimate the probability density function of a random variable based on kernels as weights
- Kernel must be chosen
  - Uniform, normal, Epanechnikov, ...
- Smoothing parameter must be chosen
  - Bias-variance trade off



# Example - Iris Flower Dataset

```
# Load iris dataset
iris = datasets.load_iris()

# Select petal length (index 2) and petal width (index 3)
X = iris.data[:, [2, 3]]
y = iris.target

# Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)
```

**iris setosa**



petal    sepal

**iris versicolor**



petal    sepal

**iris virginica**



petal    sepal

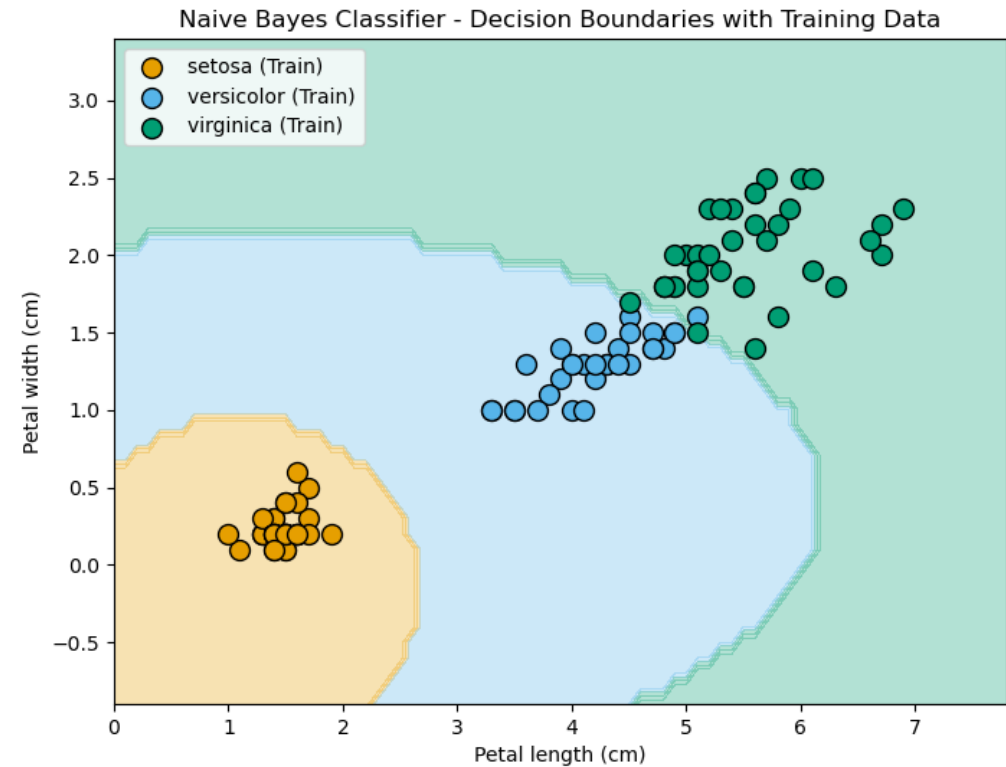
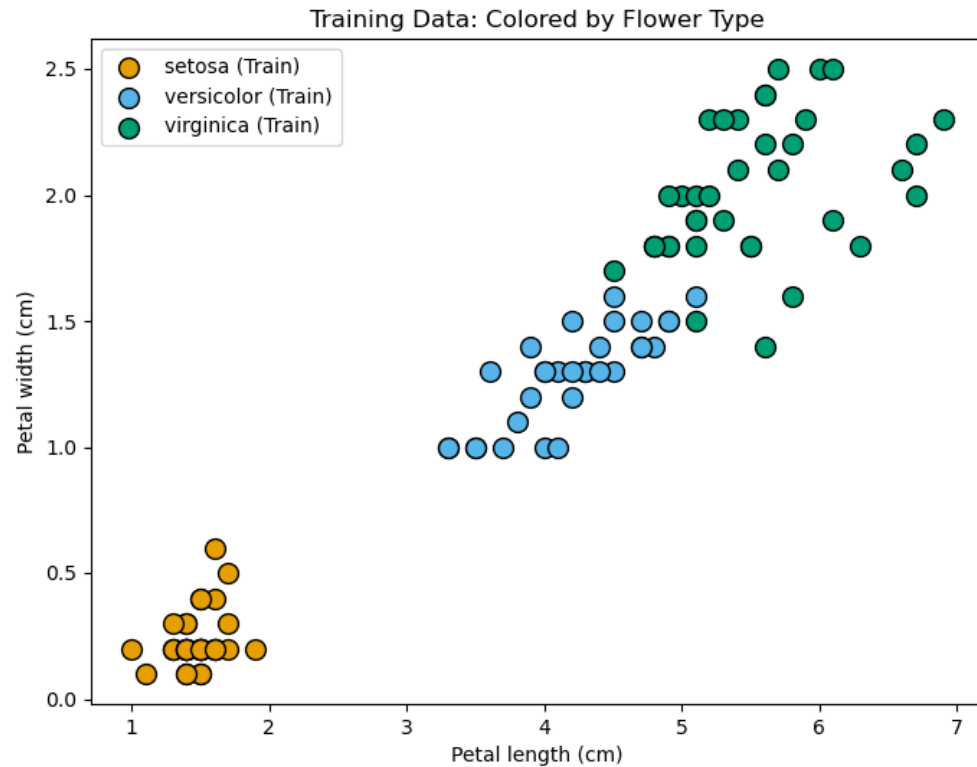
# Example – Iris Flower Dataset

```
# Training the model
gnb = GaussianNB()
gnb.fit(X_train, y_train)

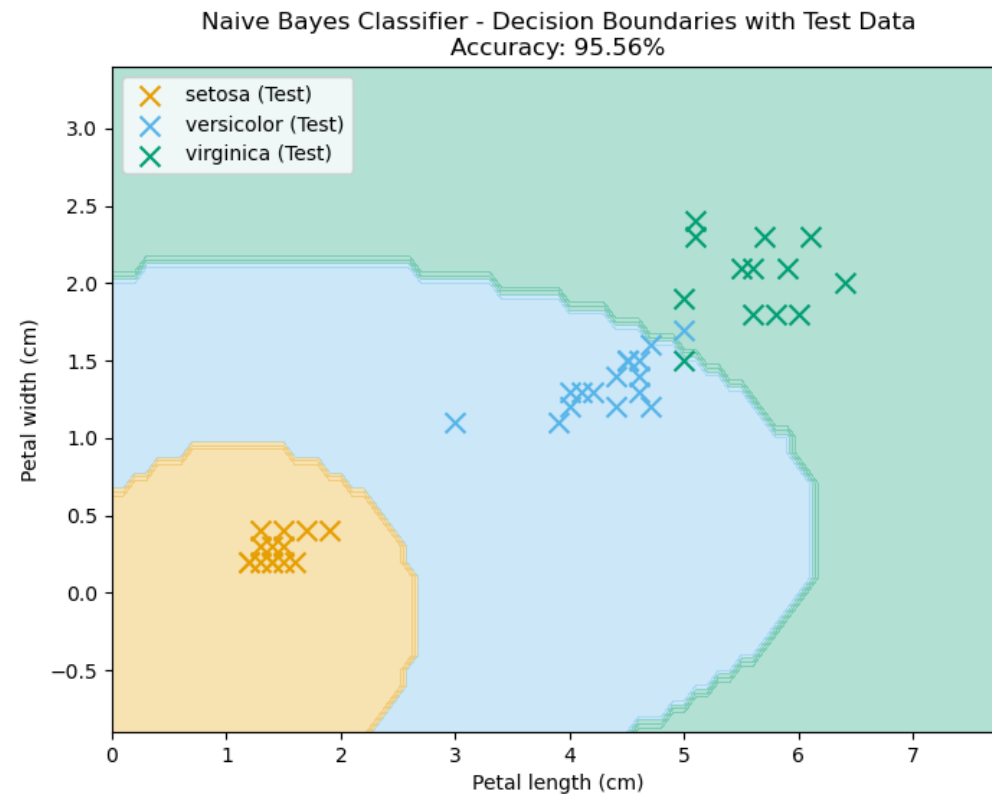
y_pred = gnb.predict(X_test) # Predicting test data labels
accuracy = metrics.accuracy_score(y_test, y_pred) * 100 # Calculate accuracy
num_train_samples = len(y_train)
num_test_samples = len(y_test)
num_correct_predictions = (y_pred == y_test).sum()
```

- Number of training samples: 105
- Number of test samples: 45
- Number of correct predictions: 43
- Accuracy: 95.56%

# Example - Iris Flower Dataset



# Example - Iris Flower Dataset






# Example - Spam Email Detection

```
[ ] # Load the dataset (replace 'email.csv' with the correct path if necessary)
emails_df = pd.read_csv('email.csv')
```

```
[ ] # Preprocessing the data
# Convert 'Category' column to binary labels (ham = 0, spam = 1)
emails_df['Category'] = emails_df['Category'].map({'ham': 0, 'spam': 1})
```

```
[ ] # Drop rows with missing values
emails_df_clean = emails_df.dropna()
```

```
[ ] emails_df_clean.head(5)
```



	Category	Message
0	0.0	Go until jurong point, crazy.. Available only ...
1	0.0	Ok lar... Joking wif u oni...
2	1.0	Free entry in 2 a wkly comp to win FA Cup fina...
3	0.0	U dun say so early hor... U c already then say...
4	0.0	Nah I don't think he goes to usf, he lives aro...

# Example - Spam Email Detection

```
[9] # Feature extraction (convert the text into numerical data using Bag-of-Words)
vectorizer = CountVectorizer()
X = vectorizer.fit_transform(emails_df_clean['Message']) # Feature matrix
y = emails_df_clean['Category'] # Labels
```

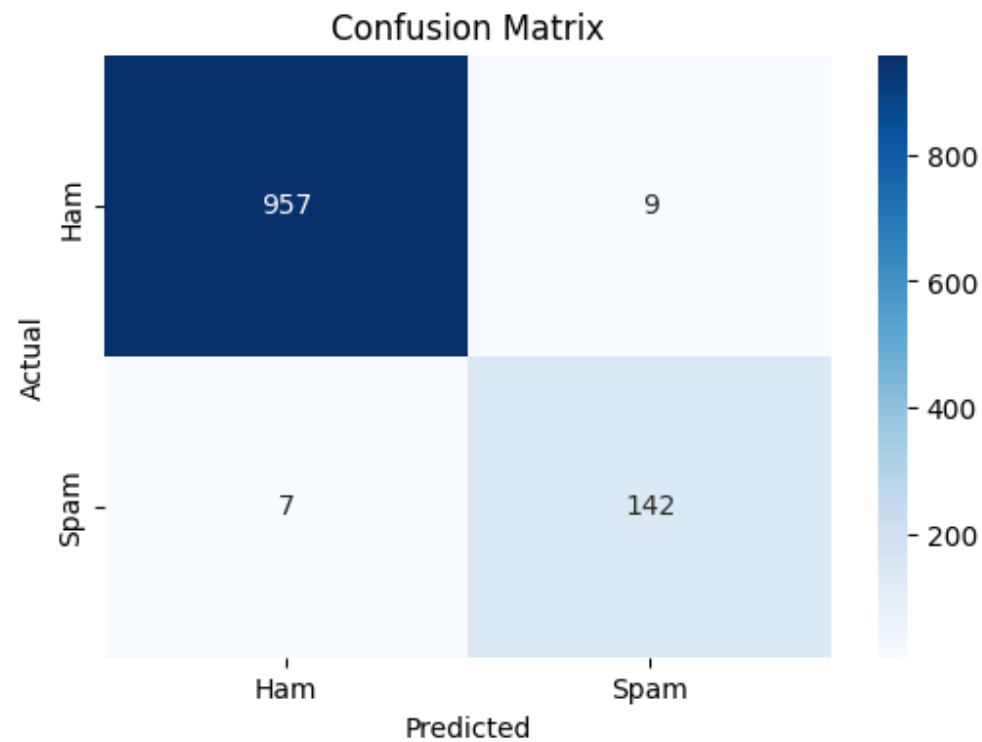
```
[10] # Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
[11] # Train the Naive Bayes classifier
nb_classifier = MultinomialNB()
nb_classifier.fit(X_train, y_train)
```

↔ MultinomialNB ⓘ ⓘ  
MultinomialNB()

- Number of training samples: 4458
- Number of test samples: 1114

# Example - Spam Email Detection



Accuracy: 0.9857

Classification Report:

	precision	recall	f1-score	support
0.0	0.99	0.99	0.99	966
1.0	0.94	0.95	0.95	149
accuracy			0.99	1115
macro avg	0.97	0.97	0.97	1115
weighted avg	0.99	0.99	0.99	1115

# Example - Spam Email Detection

```
✓ [17] # Test with a new email
0s new_email = ["We would like to congratulate you for your new position as our PhD Candidate at NTNU."]
new_email_transformed = vectorizer.transform(new_email)
prediction = nb_classifier.predict(new_email_transformed)

print(f"\nNew email prediction: {'Spam' if prediction[0] == 1 else 'Not Spam'}")
```



New email prediction: Not Spam

```
✓ [18] # Test with a new email
0s new_email = ["Congratulation! You just won a million US dollar, click here for claim."]
new_email_transformed = vectorizer.transform(new_email)
prediction = nb_classifier.predict(new_email_transformed)

print(f"\nNew email prediction: {'Spam' if prediction[0] == 1 else 'Not Spam'}")
```



New email prediction: Spam



# Common cases:

- Sentiment analysis
- Credit scoring risk prediction

# Pros and Cons

## Pros:

- Fast on high dim. data
- Good performance on limited training data

## Cons:

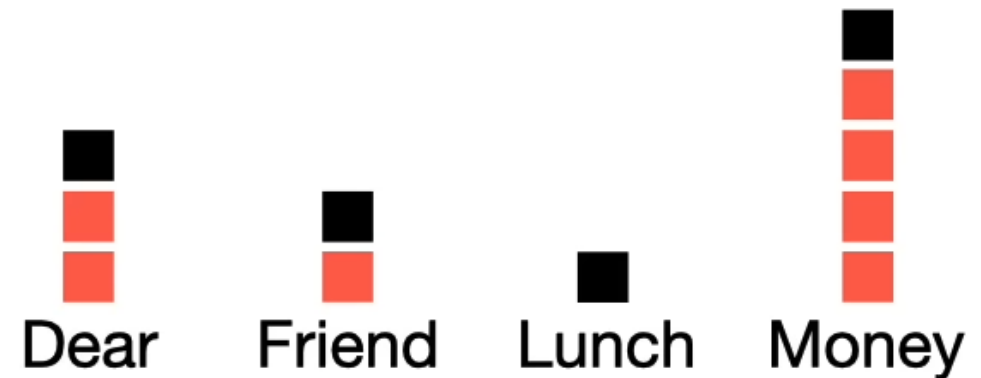
- Feature independence
- Diff. Problems with diff. versions

# What Not To Do

- Not eliminating correlated features
  - Ignoring independence assumption
  - Result in overestimation/inaccurate classification
  - Solution: DR and feature selection

# What Not To Do

- Zero frequency problem
  - Class that does not exist in training data
  - Solution: additive smoothing





# What Not To Do

- Imbalanced classes
  - Solution: resampling or adjusting class weights



# Sources

- <https://www.ibm.com/topics/naive-bayes>